

Practice with basic bioinformatics

Introduction

In order to get some hands-on experience with common practices in bioinformatics, we will work through a guided example in class today. We have four goals for today:

1. Cover the basics of how to use the **Muscle** aligner and **hmm3**.
2. Discuss some simple principles of how the quality of sequence matches are quantified.
3. Review how to generate and edit tabular information with Unix, load it into **Python/R**, and make use of it once we have it loaded into our scripting language.
4. Work through as many challenges as we have time.

Basics of Muscle

Recall that **Muscle** is designed to generate multiple sequence alignments. These alignments attempt to identify and “line up” conserved residues in sequences. Once an alignment is generated it can be used to infer evolutionary relationships as the basis of a molecular-based phylogenetic tree or create profile hidden markov models (HMMs).

In **Unix** or **Cygwin** call the **Muscle** function with the **-h** flag. Based on this help page, what is the basic syntax for generating a sequence alignment? What is the default format for an input sequence file?

Basics of hmm3

Remember that **hmm3** is used to find sequences that match a particular pattern of interest. We use **hmm3** to describe a pattern of interest with a profile HMM. A common application of **hmm3** is to generate a profile HMM for a particular protein family, domain, or active site sequence. We’ll practice that very skill today using bacterial proteomes and reference sequences from a database called *pfam* (<http://pfam.xfam.org/>). The two key **hmm3** tools we’ll need are **hmm3build** (to build a profile HMM from a sequence alignment) and **hmm3search** (to search a sequence database with a profile HMM).

In **Unix** or **Cygwin** call **hmm3build** and **hmm3search** with the **-h** flag. Based on those help pages, what is the basic syntax for generating a profile HMM and searching a database with a profile HMM?

Interpreting results

Commonly, any sequence comparison or sequence database searching tool will return values that describe the quality of the match between two sequences or between a profile HMM and a sequence. Two very common values are an **e-value** and a **bit score**. Both of these are based on an **alignment score**.

1. **alignment score** is a value that describes the quality of a pairwise or multi-sequence alignment. An exact match at a position in the alignment adds to the score and mismatches deduct from the score. Often some substitutions incur larger deductions than others based on biological principles. In addition, deductions or penalties are often included for insertion of gaps. The score for the whole alignment is the sum of the additions or deductions at each position.
2. **E-value** is a parameter that describes the number of hits one can expect to see by chance when searching a database of a particular size. It decreases exponentially as the *alignment score* of the match increases. This means that the *E-value* is sensitive to the size of the database, the score of the alignment,

and the length of the alignment. The closer the E-value is to zero, the more “significant” the match is. Keep in mind that even identical short alignments have relatively high E-values because shorter sequences have a higher probability of occurring in the database purely by chance.

3. **bit score** is the required size of a database in which the current match could be found just by chance. It is presented on a \log_2 scale. This means each increase by one doubles the required database size. This has become a more popular metric as it does not depend on database size, like *E-value*.

What is the E-value or bit score that indicates a “true” match?

Processing tabular information

Unix: Three Unix tools are useful for editing tabular files.

1. **grep** is useful for getting rid of header lines. Nice bioinformatics tools will indicate header lines with some special character. For example, **hmmsearch** header lines have ‘#’ at the beginning. Remember the **-v** argument will return all lines that DO NOT match your search string.
2. **cut** can be used to capture specific columns from tabular data. Remember it takes a **-d** argument to set the delimiter between columns and a **-f** to specify which columns to return.
3. **awk** is a very powerful tool in Unix, but we will focus on one specific functionality today. For example, **awk '{print \$1,\$2}'** would print the first and second column of a file separated by a space.

*How could we reduce tabular output to the 1st, 4th, and 8th column of a tabular text file returned by **hmmsearch**?*

Python/R: Don’t forget also that we can use dataframes and vectors in **R** or dataframes, lists, and dictionaries in **Python** to hold data, which can then be manipulated using subsetting techniques we’ve practiced over the last few weeks.

*How could we find information contained in one table that is associated with sequence identifiers that were identified as sequences of interest by **hmmsearch**?*

Challenges

On Sakai, there are eight fasta files that each contain the protein sequences of genes in a bacterial genome, an additional fasta file with reference sequences for a domain of a gene encoding a gene expression regulator called sigma70, and a file containing annotations for each protein sequence in the *Roseobacter* genome. Download these files and work through the challenges below.

- 1) Align the sequences in sigma70.fasta using muscle. How does the alignment file differ from the original file?
- 2) Build a profile HMM using the sigma70 sequence alignment and hmmbuild.
- 3) Search the *Roseobacter* proteome (**Roseobacter.fasta**) for sigma factors using hmmsearch. How many matches do you get?

I suggest using the argument for tabular output, as this will be useful in other steps.

- 4) Make a histogram of the bit scores for all sigma factor hits from the *Roseobacter* proteome.

*I would use **Python/R** for this after creating an easy to load table in **Unix**.*

- 5) Use the annotations of genes from the *Roseobacter* proteome (**Roseobacter.annot**) to check whether your sigma factor hits seem like good matches.

*Again, I would use **Python/R** for this after creating an easy to load table in **Unix**.*

- 6) How many proteins are encoded in each of the eight proteomes provided?

*Remember how to use a for loop in **Unix**?*